


Operating System Principles

操作系统原理



File System

李旭东

leexudong@nankai.edu.cn

Nankai University

Long-Term Information Storage

- Essential requirements
 - It must be possible to store a very large amount of information
 - The information must survive the termination of the process using it
 - Multiple processes must be able to access the information concurrently
- Solution
 - ???

Long-Term Information Storage

- Solution

- Store information on disk and other external media in units called files.
- Processes can then read them and write new one if need be
- Persistent
- File, Directory, File System


Long-Term Information Storage

- Storage with Block Device

- Magnetic Disks
- Magnetic Tape
- Optical disc
- Flash memory

- Block-Level Operations

- Read i -th Block
- Write i -th Block



Long-Term Information Storage

- Retrieve Information Quickly
 - 1. How do you find information?
 - 2. How do you keep one user from reading another user's data?
 - 3. How do you know which blocks are free?



Objectives

- File Concept
- Directory Concept
- File Share & Protection
- File System Implementation
- File System Reliability
- File System Performance
- File System Cases



File Concept

- File
 - A logical units of information
 - A byte stream
- File Size
- File Name
- File Logical Structure (File Content)
- File Type
- File Access
- File Attributes
- File Operations
- File Physical Structure



File Size

- Zero Byte
- ~~1 Bit~~
- 1 Byte
- 2 Bytes
- 1 KB (千)
- 1 MB (兆)
- ...
- GB (吉)
- TB (太)
- PB (拍)
- EB (艾)
- ZB (泽)
- 1.2ZB
- YB (尧)
- ...

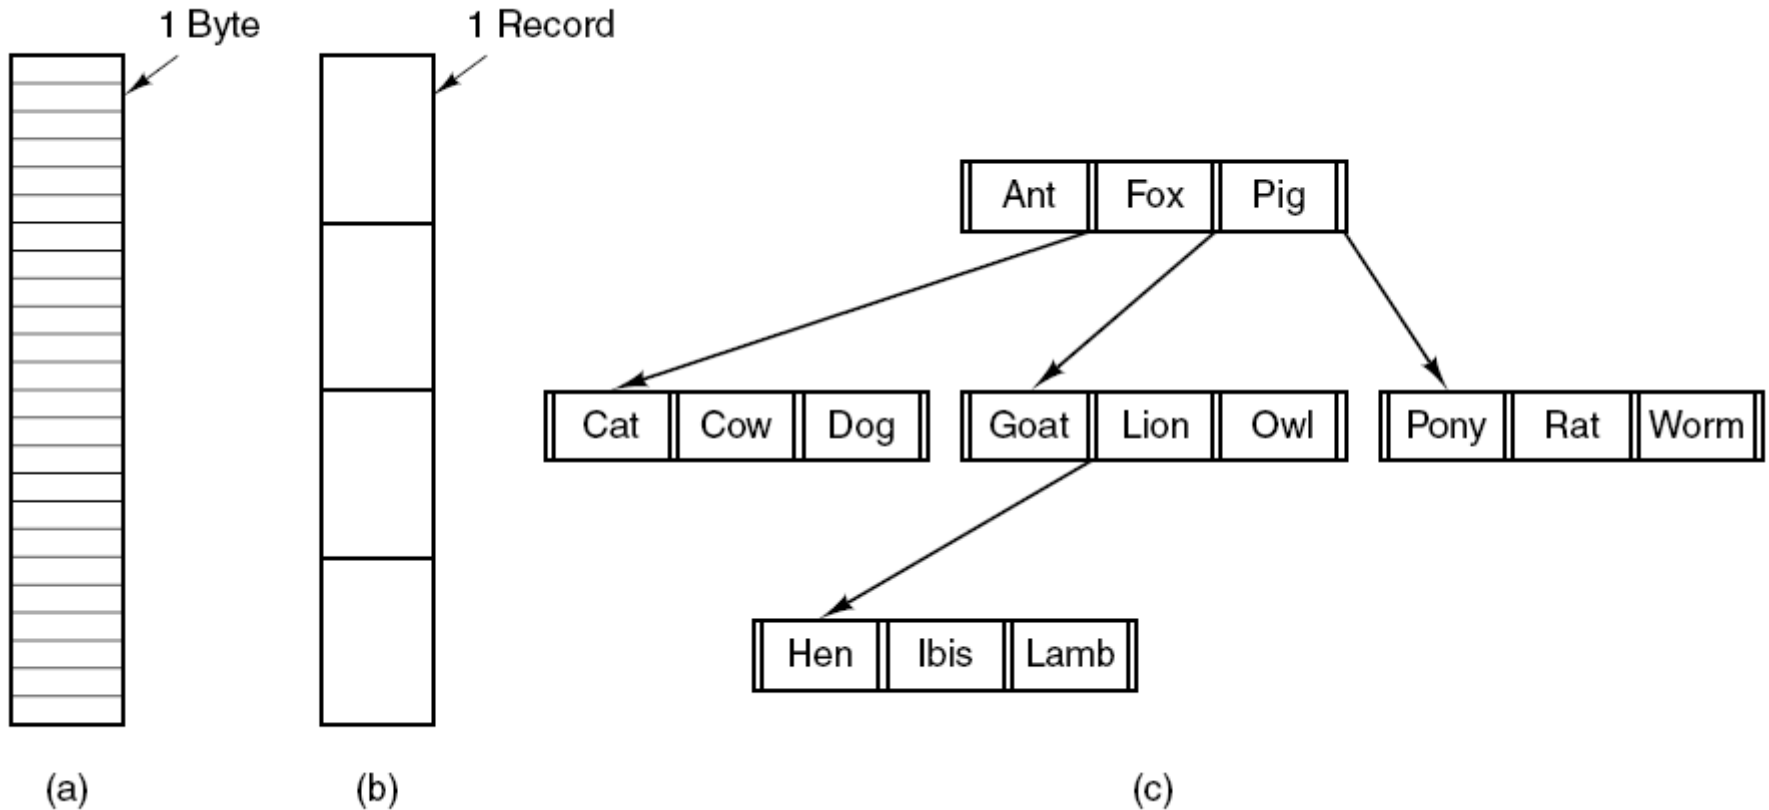


File Naming

- File extension
- File Name Charset

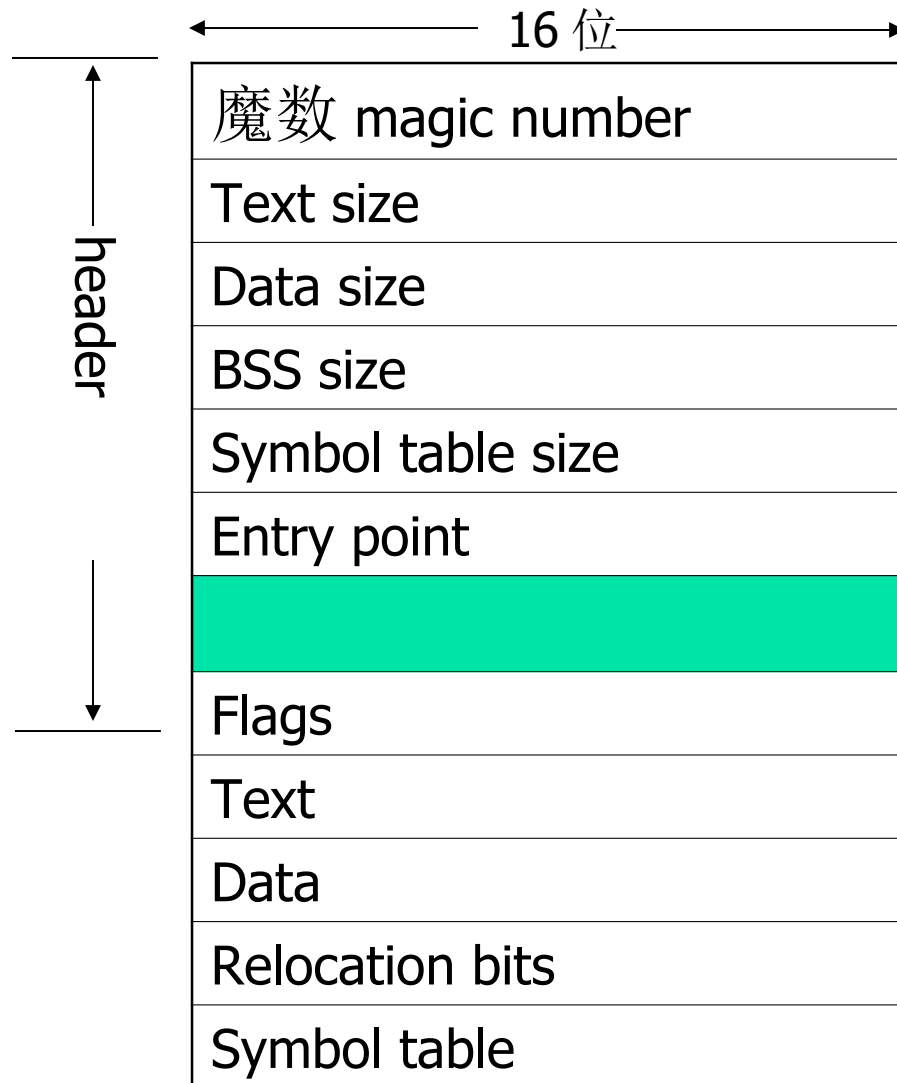
Extension	Meaning
file.bak	Backup file
file.c	C source program
file.gif	Compuserve Graphical Interchange Format image
file.hlp	Help file
file.html	World Wide Web HyperText Markup Language document
file.jpg	Still picture encoded with the JPEG standard
file.mp3	Music encoded in MPEG layer 3 audio format
file.mpg	Movie encoded with the MPEG standard
file.o	Object file (compiler output, not yet linked)
file.pdf	Portable Document Format file
file.ps	PostScript file
file.tex	Input for the TEX formatting program
file.txt	General text file
file.zip	Compressed archive

File Logical Structure

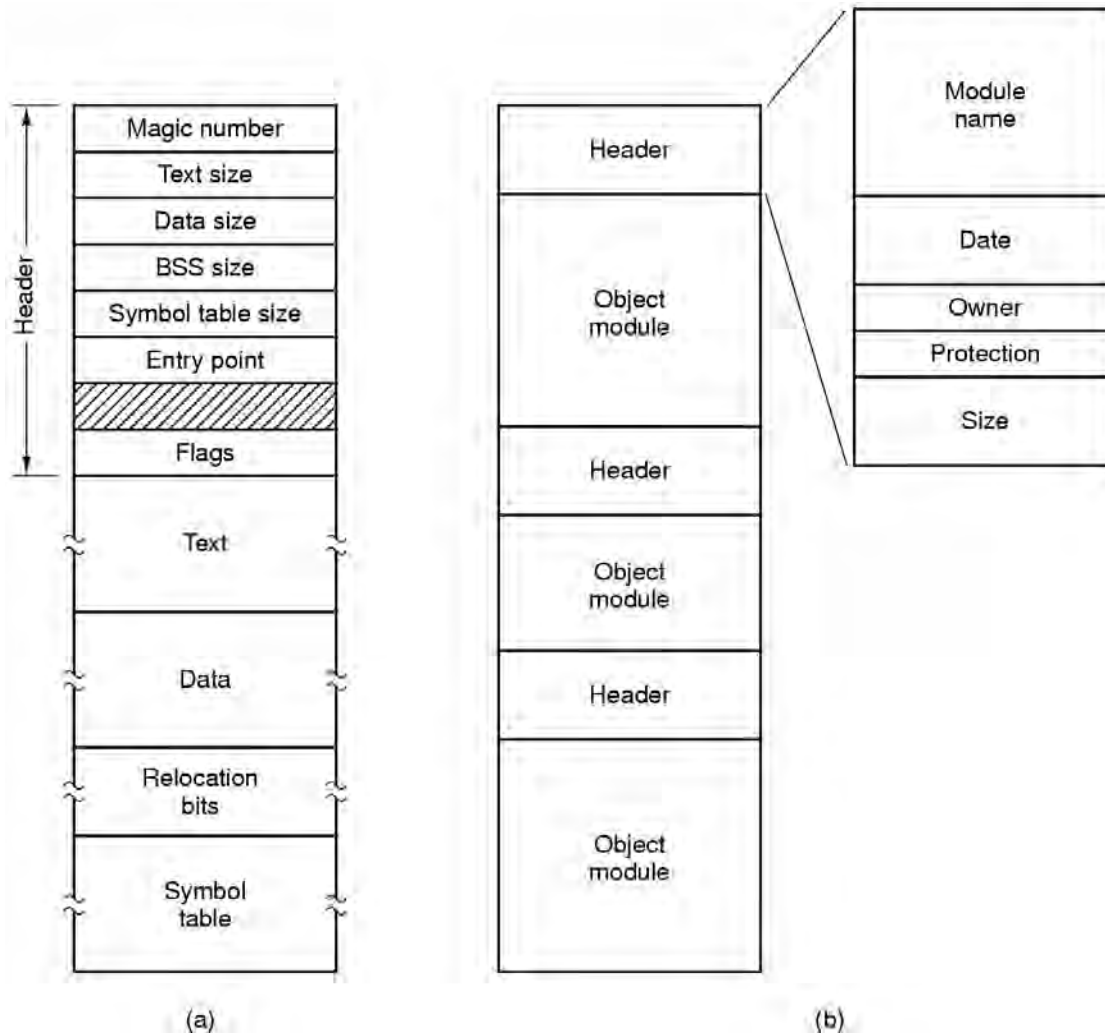


(a) Byte sequence. (b) Record sequence. (c) Tree.

File Logical Structure: Executable Program File



File Types



(a) An executable file. (b) An archive.



File Types

- T1
 - Regular files
 - Directory files
- T2
 - Character special files
 - Block speicial files
- T3
 - ASCII files
 - Binary files
- ...



File Access

- Sequential access
- Random access
- Key-value access



File Attributes

Attribute	Meaning
Protection	Who can access the file and in what way
Password	Password needed to access the file
Creator	ID of the person who created the file
Owner	Current owner
Read-only flag	0 for read/write; 1 for read only
Hidden flag	0 for normal; 1 for do not display in listings
System flag	0 for normal files; 1 for system file
Archive flag	0 for has been backed up; 1 for needs to be backed up
ASCII/binary flag	0 for ASCII file; 1 for binary file
Random access flag	0 for sequential access only; 1 for random access
Temporary flag	0 for normal; 1 for delete file on process exit
Lock flags	0 for unlocked; nonzero for locked
Record length	Number of bytes in a record
Key position	Offset of the key within each record
Key length	Number of bytes in the key field
Creation time	Date and time the file was created
Time of last access	Date and time the file was last accessed
Time of last change	Date and time the file was last changed
Current size	Number of bytes in the file
Maximum size	Number of bytes the file may grow to



File Operations

The most common system calls relating to files:

- Create
- Delete
- Open
- Close
- Read
- Write
- Append
- Seek
- Get Attributes
- Set Attributes
- Rename

CP Program:Using File System Calls (1/2)

```
/* File copy program. Error checking and reporting is minimal. */

#include <sys/types.h>                /* include necessary header files */
#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]);    /* ANSI prototype */

#define BUF_SIZE 4096                /* use a buffer size of 4096 bytes */
#define OUTPUT_MODE 0700             /* protection bits for output file */

int main(int argc, char *argv[])
{
    int in_fd, out_fd, rd_count, wt_count;
    char buffer[BUF_SIZE];

    if (argc != 3) exit(1);           /* syntax error if argc is not 3 */
```

CP Program:Using File System Calls (2/2)

```
/* Open the input file and create the output file */
in_fd = open(argv[1], O_RDONLY); /* open the source file */
if (in_fd < 0) exit(2);           /* if it cannot be opened, exit */
out_fd = creat(argv[2], OUTPUT_MODE); /* create the destination file */
if (out_fd < 0) exit(3);          /* if it cannot be created, exit */

/* Copy loop */
while (TRUE) {
    rd_count = read(in_fd, buffer, BUF_SIZE); /* read a block of data */
    if (rd_count <= 0) break;                  /* if end of file or error, exit loop */
    wt_count = write(out_fd, buffer, rd_count); /* write data */
    if (wt_count <= 0) exit(4);                 /* wt_count <= 0 is an error */
}

/* Close the files */
close(in_fd);
close(out_fd);
if (rd_count == 0) /* no error on last read */
    exit(0);
else
    exit(5); /* error on last read */
}
```



File Physical Structure

- How to store in block storage
 - ???




Directory Concept

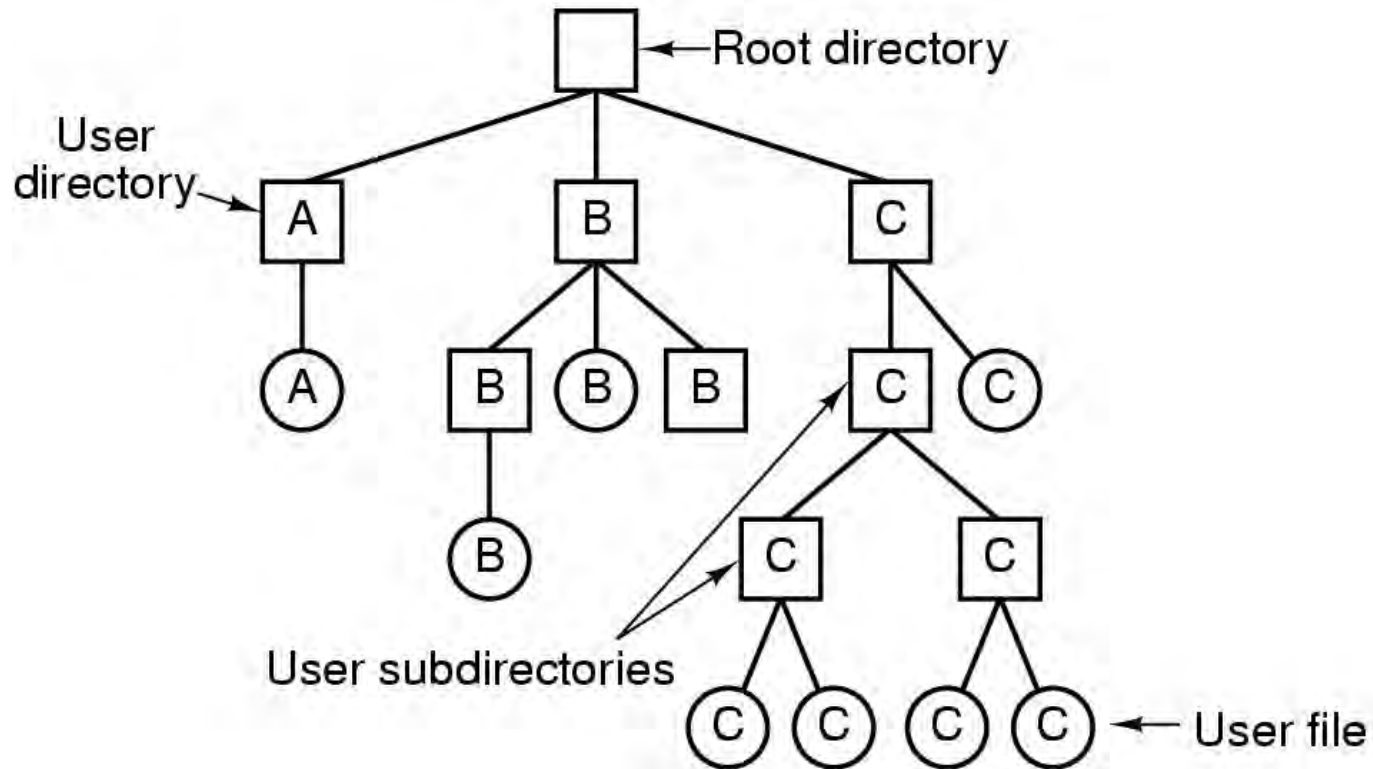
- Directory
 - Folder, Which contains files
- Single-level Directory System
- Hierarchical Directory System
- Path Names
- Directory Operations
- Implement of Directory

Single-level Directory System

Root directory

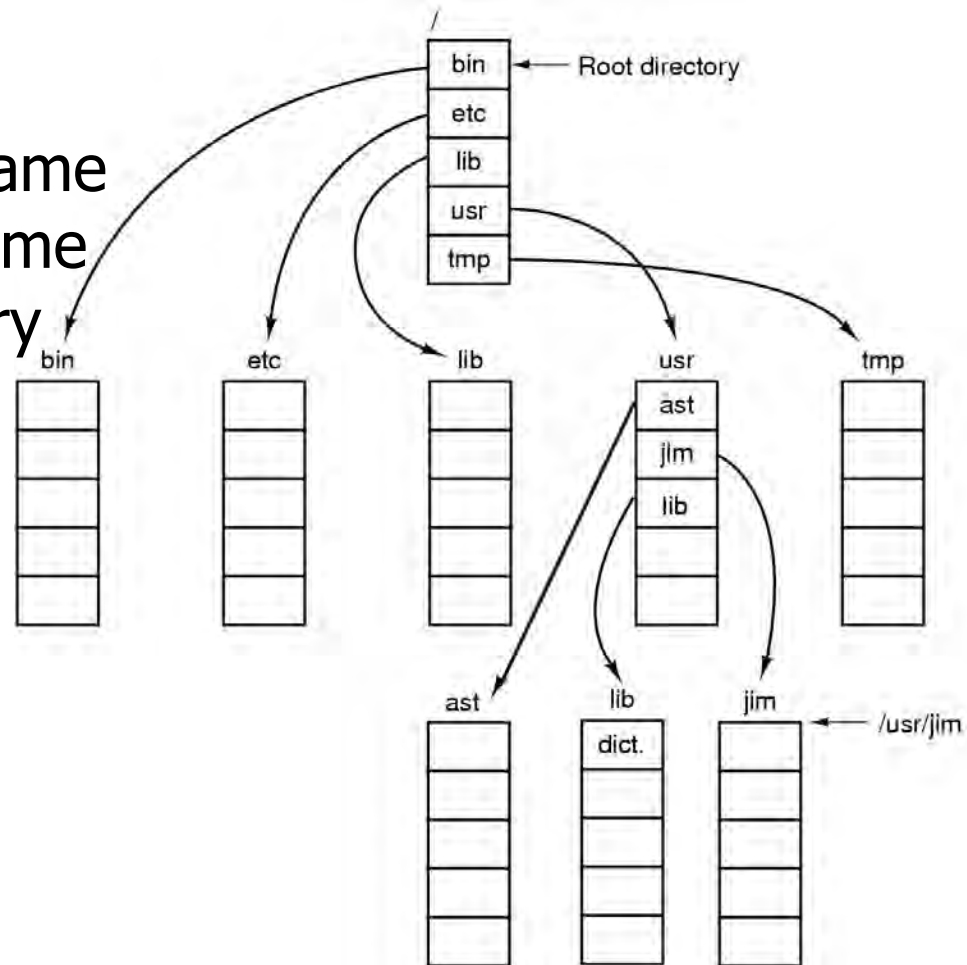


Hierarchical Directory System



Path Names

- Absolute path name
- Relative path name
- Working directory



A UNIX directory tree



Directory Operations

System calls for managing directories:

- Create
 - Delete
 - Opendir
 - Closedir
 - Readdir
 - Rename
 - Link
 - Unlink
-
- Link
 - Hard link
 - Symbolic link



Implement of Directory

- How to store in block storage
 - ???



File System Concept

- File System
 - It contains files and directories
- File System Types
- Implement of File System



File System Types

- Traditional: ext2
- Newest: ext3, ReiserFS, IBM JFS, xfs
- Other UNIX: minix, ext, xiafs
- FAT-12, FAT-16, FAT-32, VFAT, NTFS (read-only)
- HPFS (OS/2) read-only, HFS (Macintosh) read-only
- AFFS (Amiga), System V, Coherent, Xenix
- CD-ROM: **ISO 9660**
- UMSDOS (UNIX-like FS on MS-DOS)
- NFS (Network File System)
- SMBFS (Windows share), NCPFS (Novell Netware share)
- /proc (for kernel and process information)
- SHMFS (Shared Memory Filesystem)

File Systems Supported by Linux

9p	configfs	freevxfs	jffs2	nls	smbfs
adfs	cramfs	fscache	jfs	notify	squashfs
affs	debugfs	fuse	Kconfig	ntfs	sysfs
afs	devpts	gfs2	Kconfig.binfmt	ocfs2	sysv
autofs	dlm	hfs	lockd	omfs	ubifs
autofs4	ecryptfs	hfsplus	logfs	openpromfs	udf
befs	efs	hostfs	Makefile	partitions	ufs
bfs	exofs	hpfs	minix	proc	xfs
btrfs	exportfs	hppfs	ncpfs	qnx4	
cachefiles	ext2	hugetlbfs	nfs	quota	
ceph	ext3	isofs	nfs_common	ramfs	
cifs	ext4	jbd	nfsd	reiserfs	
coda	fat	jbd2	nilfs2	romfs	

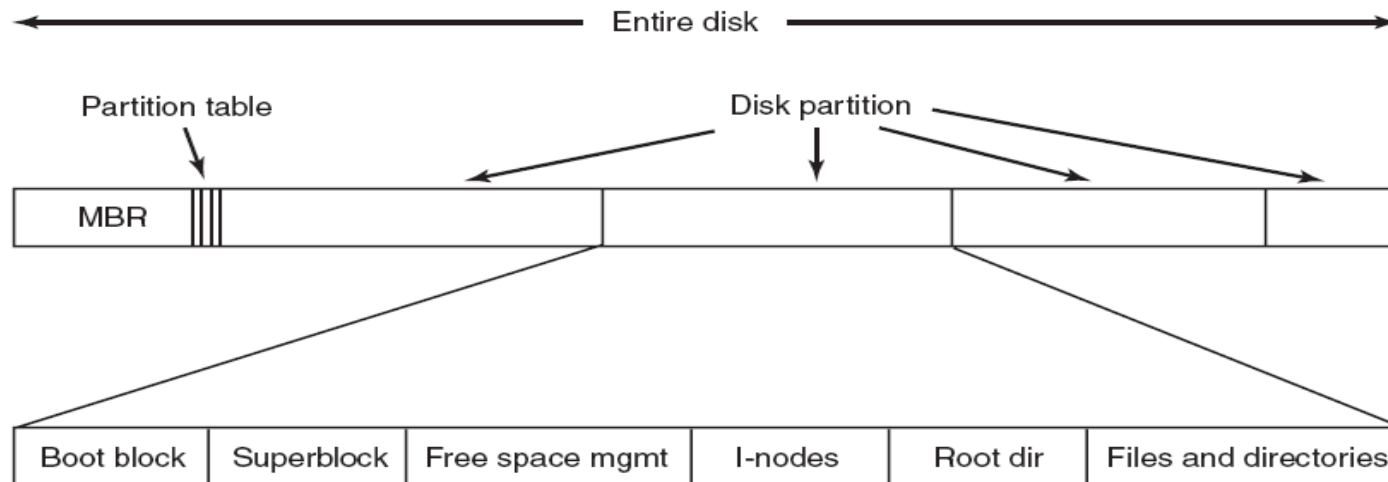


Implement of File System

- File System Layout
- Implement of Files
- Implement of Directories
- Free Block Space Management


File System Layout

- Disk
- MBR (Master Boot Record)
 - Boot block
- Partitions
 - Superblock





Implement of Files

Where are File's Blocks?



Implement of Files

■ Contiguous Allocation



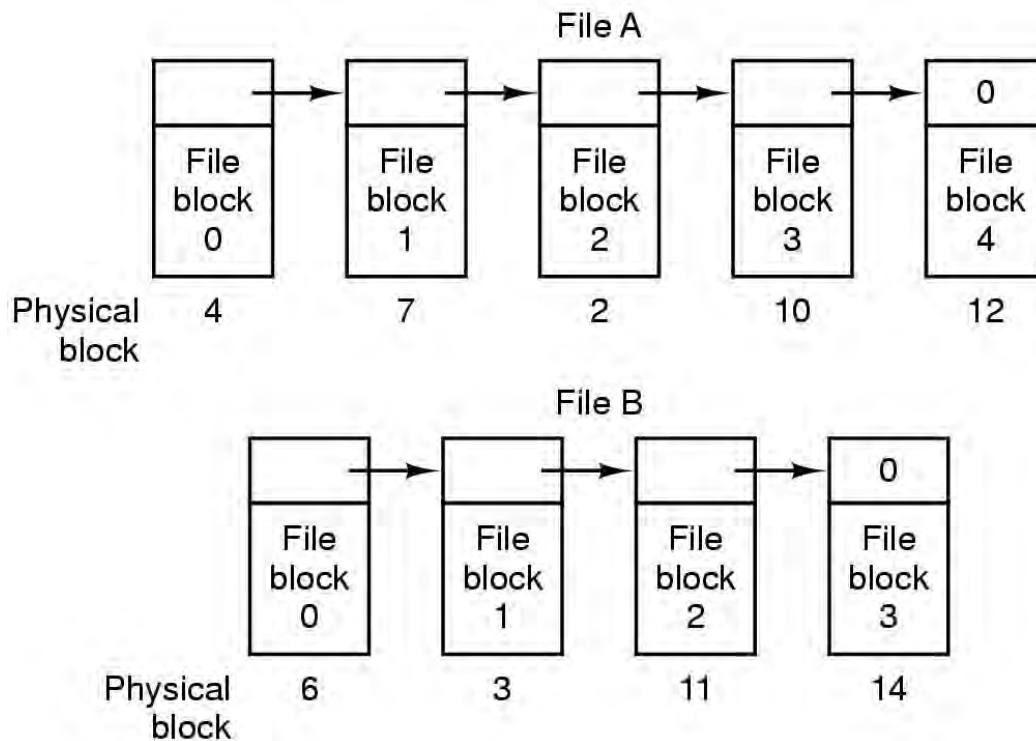
(a) Contiguous allocation of disk space for 7 files.

(b) The state of the disk after files D and F have been removed.

Implement of Files

■ Linked List Allocation

- The first word of each block is used as a pointer to the next one
- The rest of the block is for data



- Internal fragmentation
- Random access is extremely slow

Storing a file as a linked list of disk blocks.

Implement of Files

- Linked List Allocation Using a Table in Memory

- FAT: File Allocation Table

- FAT12

- FAT16

- FAT32

- Random access is much easier

- ?Disadvantage

- Too many table entries in memory

Physical
block

0	
1	
2	10
3	11
4	7
5	
6	3
7	2
8	
9	
10	12
11	14
12	-1
13	
14	-1
15	

← File A starts here

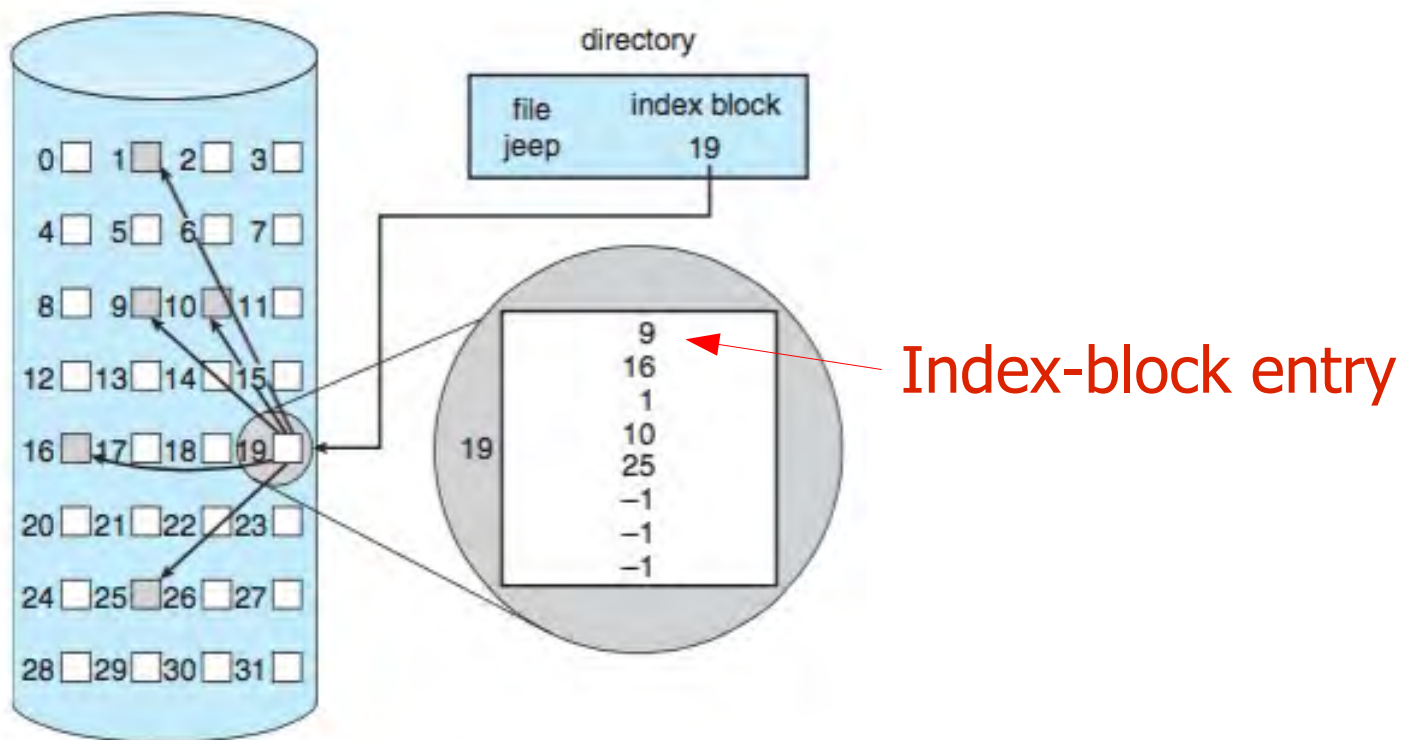
← File B starts here

← Unused block


Implement of Files

- I-nodes

- Index-node, which lists the attributes and disk addresses of the files blocks

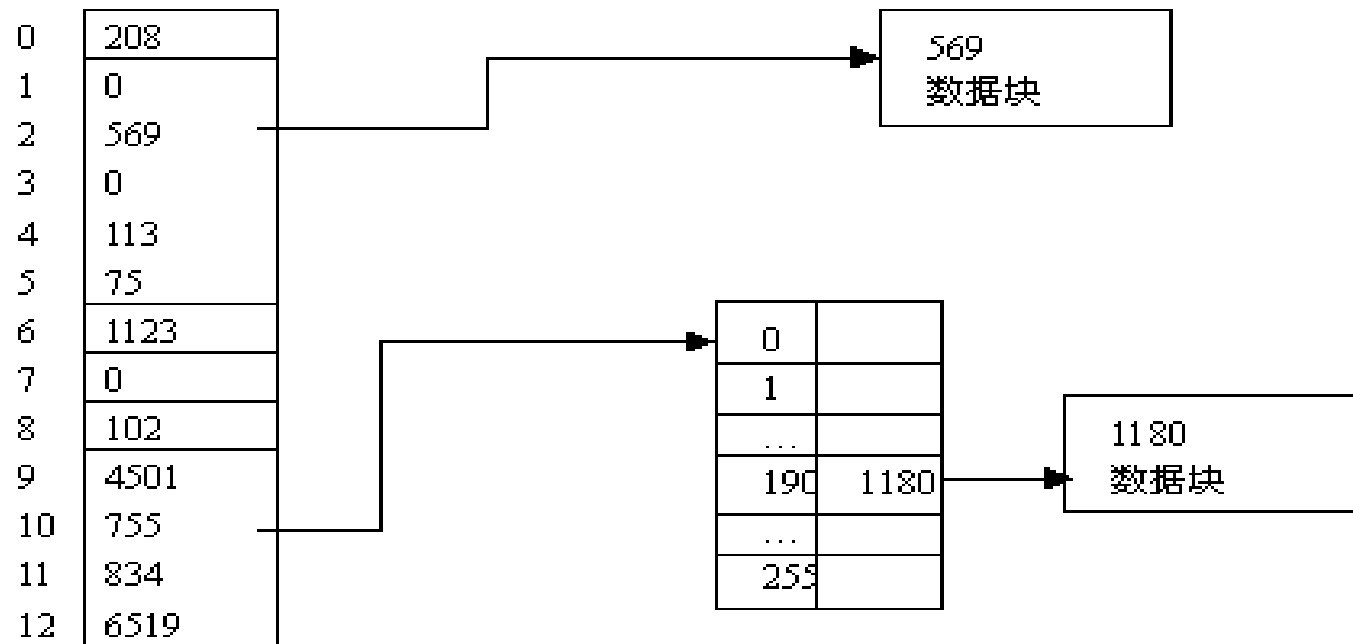


I-Nodes



UNIX Inode

Quiz



How to read the 204803th data of file A, which physical block index has 4 bytes?

Question 1. when the size of physical block is 1024?

Question 2. when the size of physical block is 2048?

Implement of Directories


- Directory entry

- File Control Block, FCB
- Fixed-size, Variable-size

- File
 - FCB
 - Body

games	attributes
mail	attributes
news	attributes
work	attributes

(a)



(b)

Data structure containing the attributes

(a) A simple directory containing fixed-size entries with the disk addresses and attributes in the directory entry.

(b) A directory in which each entry just refers to an i-node.



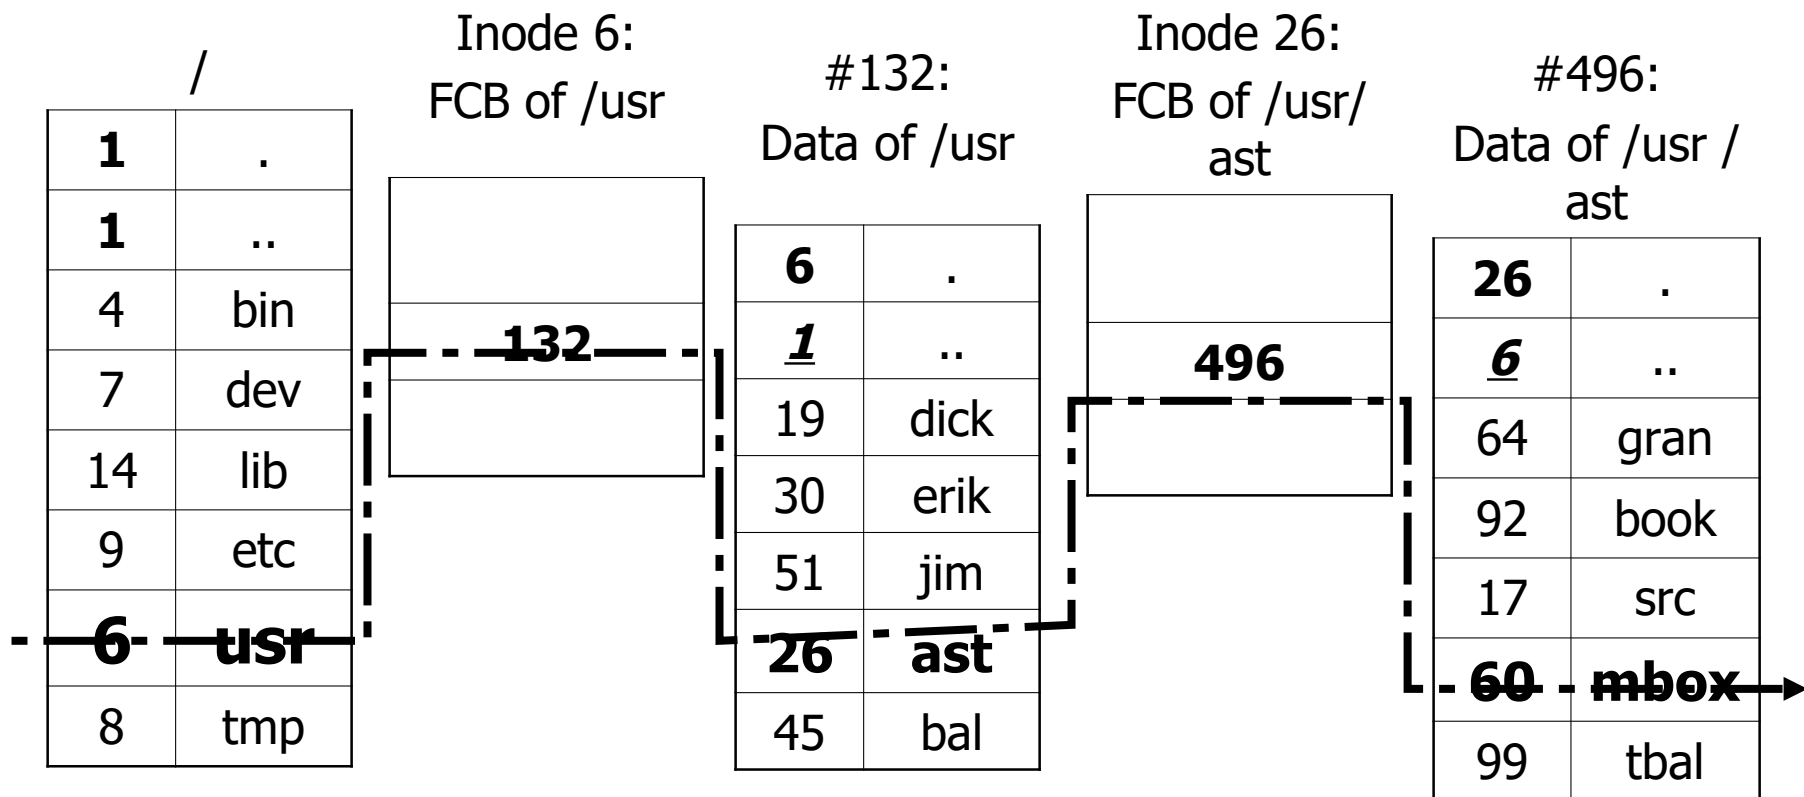
Two ways of handling long file names in a directory



(a) In-line. (b) In a heap.

Implement of Directories

- Search file by pathname
 - /usr/ast/mbox






Implement of Directories

- Directory entry
 - Linear List
 - Hash Table

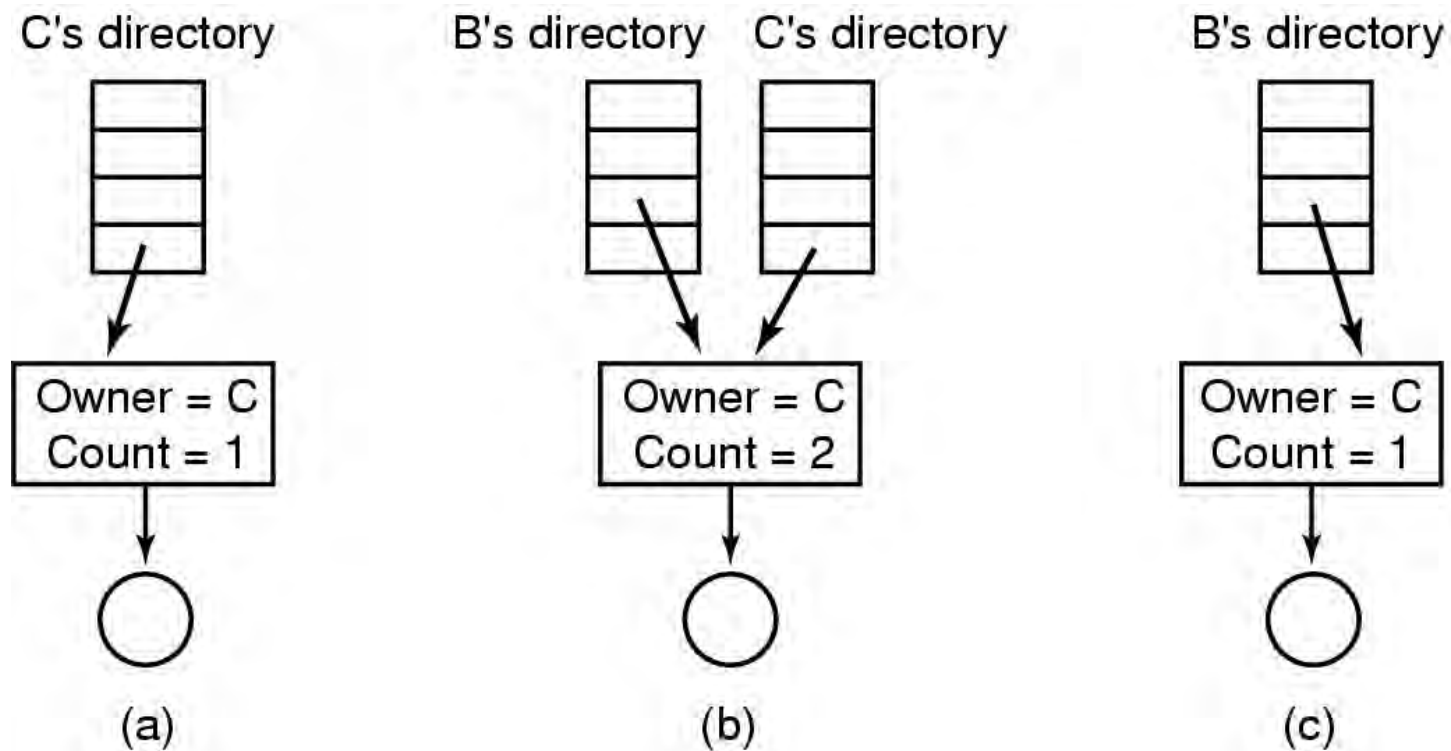
Shared Files

- Symbolic linking
- Hard linking





File system containing a shared file.

Shared Files



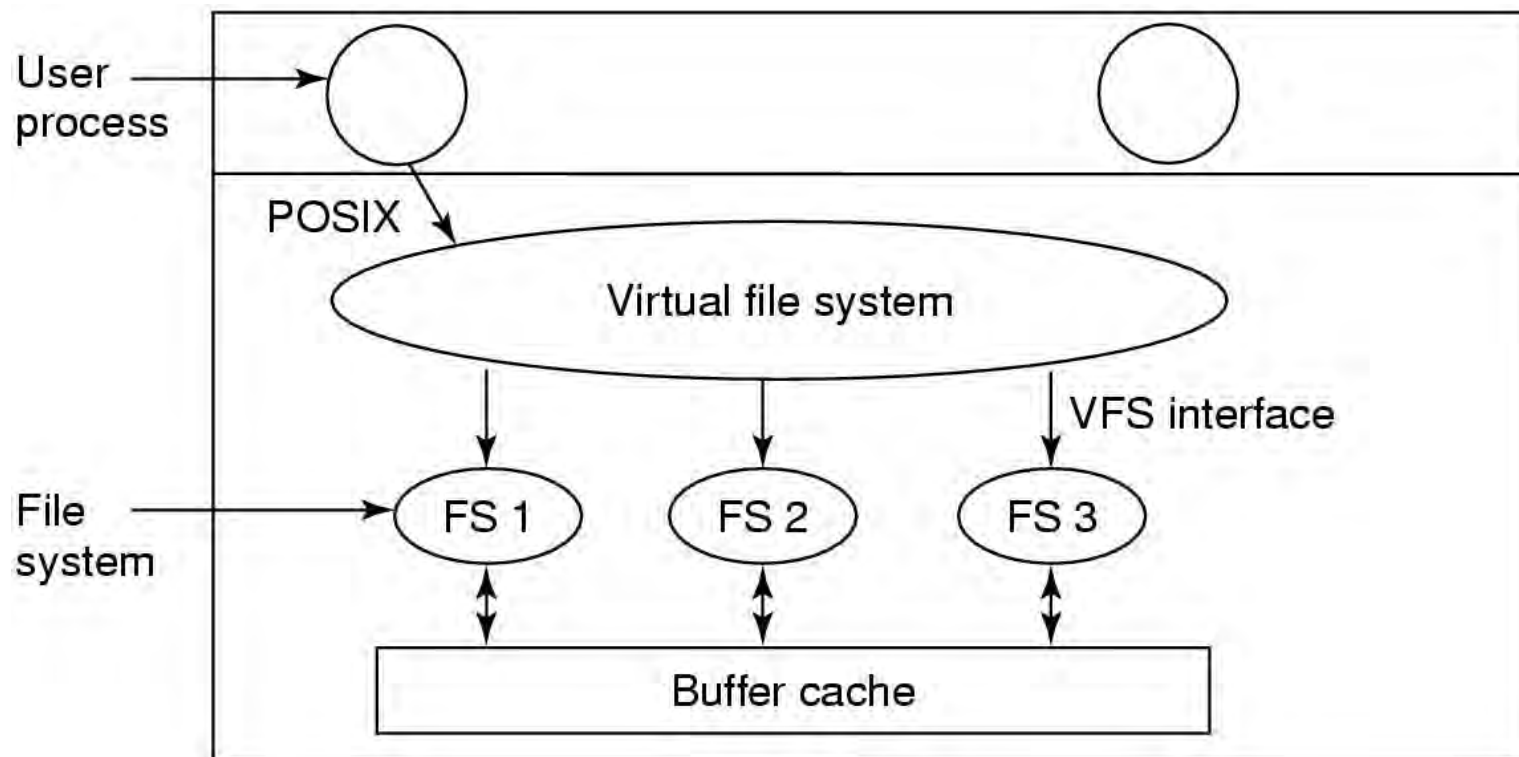
(a) Situation prior to linking. (b) After the link is created.
(c) After the original owner removes the file.

In-memory file-system structures

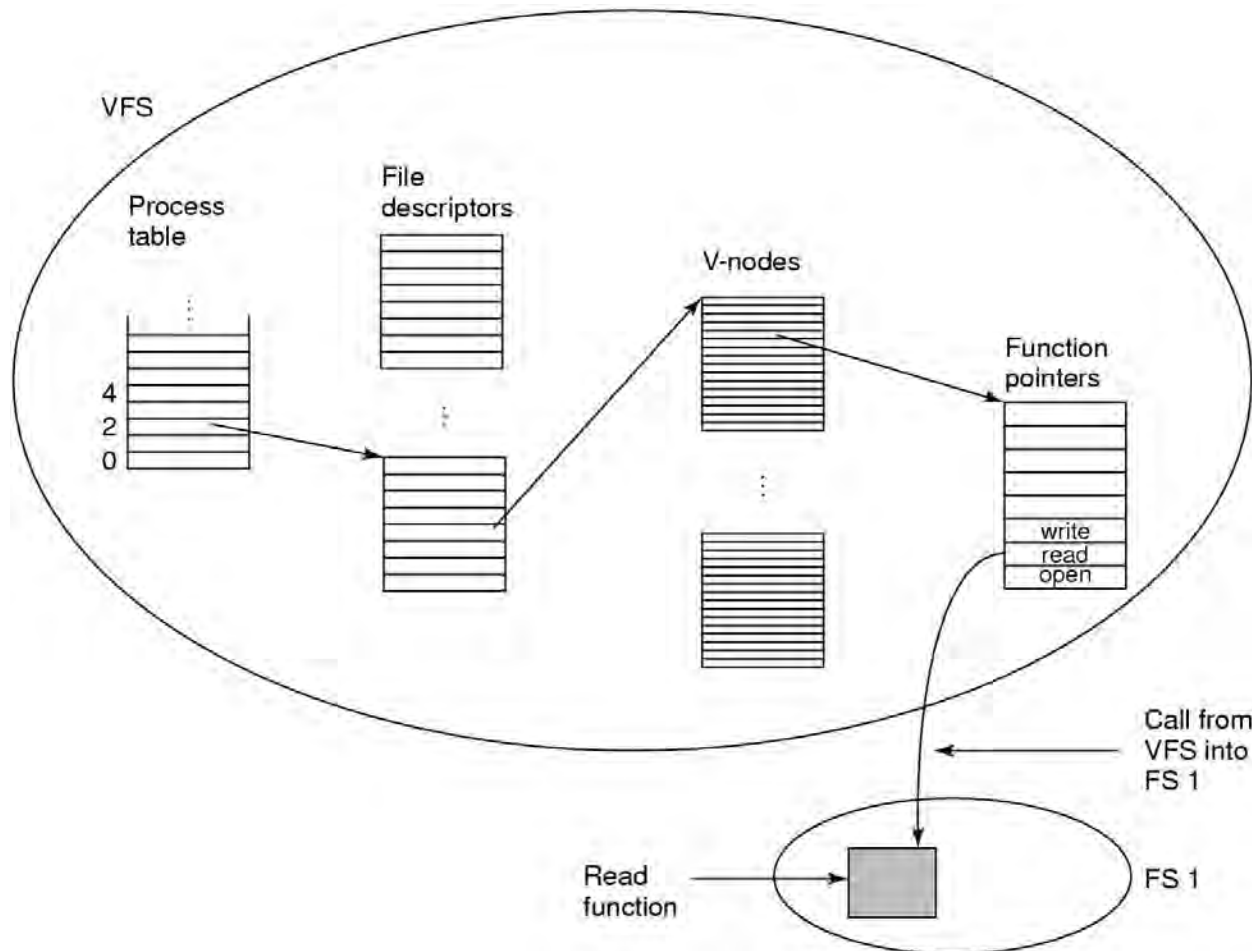


Virtual File Systems

- VFS Interface



Virtual File Systems



A simplified view of the data structures and code used by the VFS and concrete file system to do a read.



Disk Space Management

- Block Size
- Keeping Track of Free Blocks
- Disk Quotas
- File System Backups
- File System Consistency
- File System Performance
- Defragmenting Disks



Block Size

■ Small ? Large?


Length	VU 1984	VU 2005	Web
1	1.79	1.38	6.67
2	1.88	1.53	7.67
4	2.01	1.65	8.33
8	2.31	1.80	11.30
16	3.32	2.15	11.46
32	5.13	3.15	12.33
64	8.71	4.98	26.10
128	14.73	8.03	28.49
256	23.09	13.29	32.10
512	34.44	20.62	39.94
1 KB	48.05	30.91	47.82
2 KB	60.87	46.09	59.44
4 KB	75.31	59.13	70.64
8 KB	84.97	69.96	79.69

Length	VU 1984	VU 2005	Web
16 KB	92.53	78.92	86.79
32 KB	97.21	85.87	91.65
64 KB	99.18	90.84	94.80
128 KB	99.84	93.73	96.93
256 KB	99.96	96.12	98.48
512 KB	100.00	97.73	98.99
1 MB	100.00	98.87	99.62
2 MB	100.00	99.44	99.80
4 MB	100.00	99.71	99.87
8 MB	100.00	99.86	99.94
16 MB	100.00	99.94	99.97
32 MB	100.00	99.97	99.99
64 MB	100.00	99.99	99.99
128 MB	100.00	99.99	100.00



Percentage of files smaller than a given size (in bytes).

Block Size



The solid curve (left-hand scale) gives the data rate of a disk.
The dashed curve (right-hand scale) gives the disk space efficiency.
All files are 4 KB.

Keeping Track of Free Blocks


- Bit Vector
 - Bitmap

1001101101101100
0110110111110111
1010110110110110
0110110110111011
1110111011101111
1101101010001111
0000111011010111
1011101101101111
1100100011101111
≈
0111011101110111
1101111101110111

A bitmap


Keeping Track of Free Blocks

- Linked List

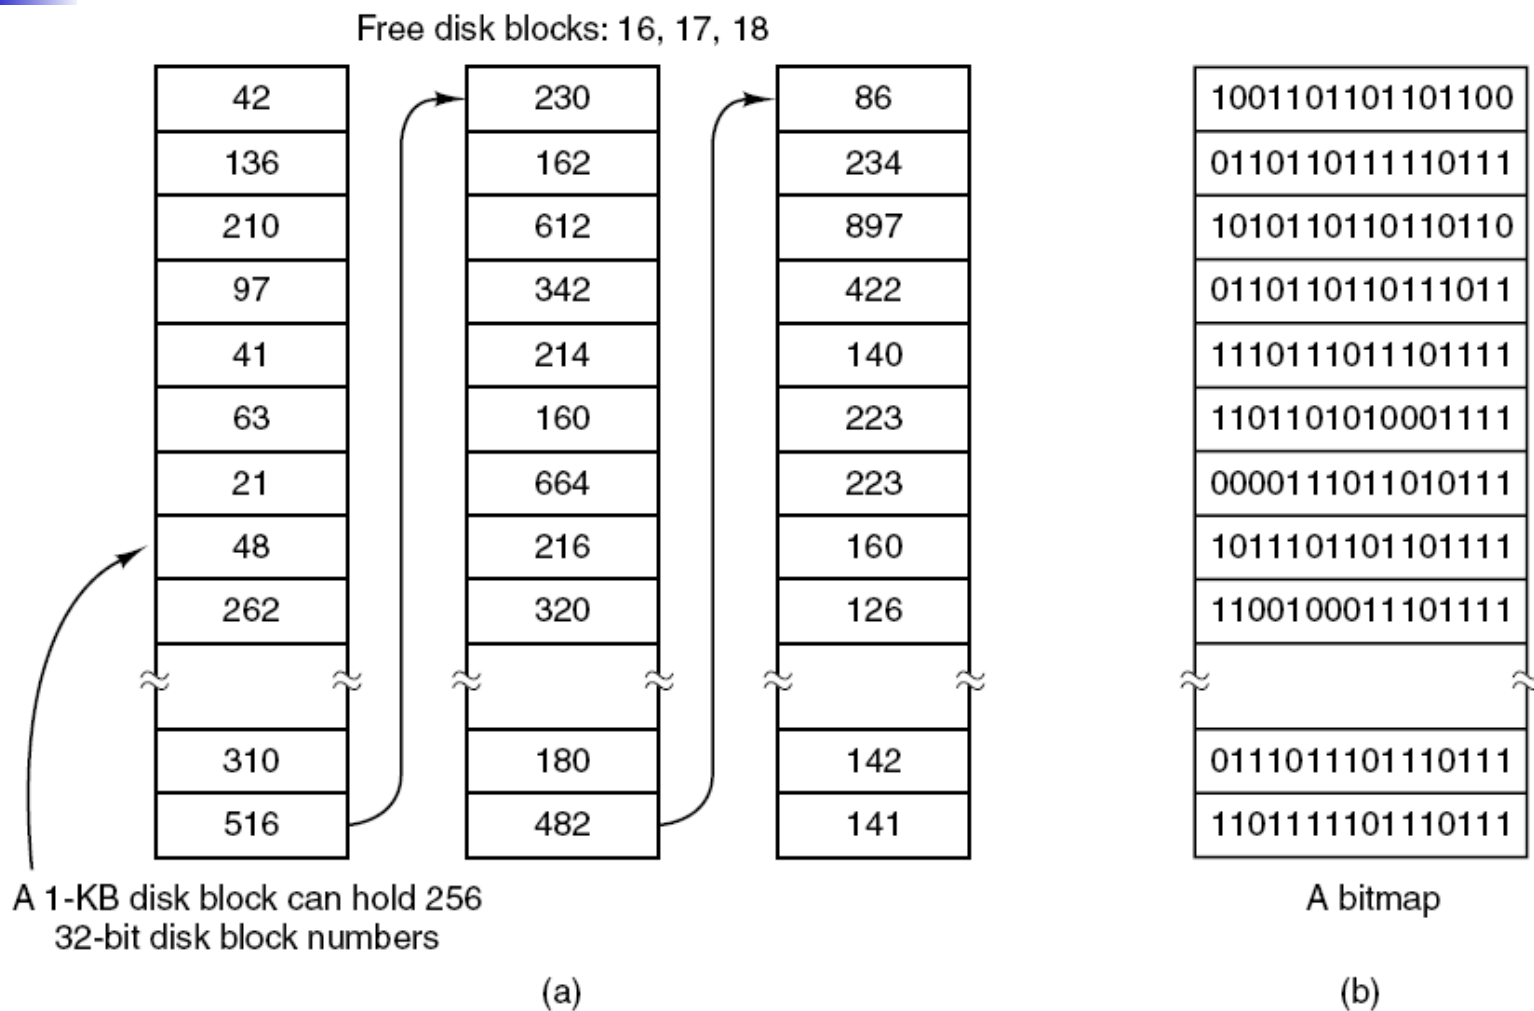


Keeping Track of Free Blocks

- Grouping Linked List: 成组连接法

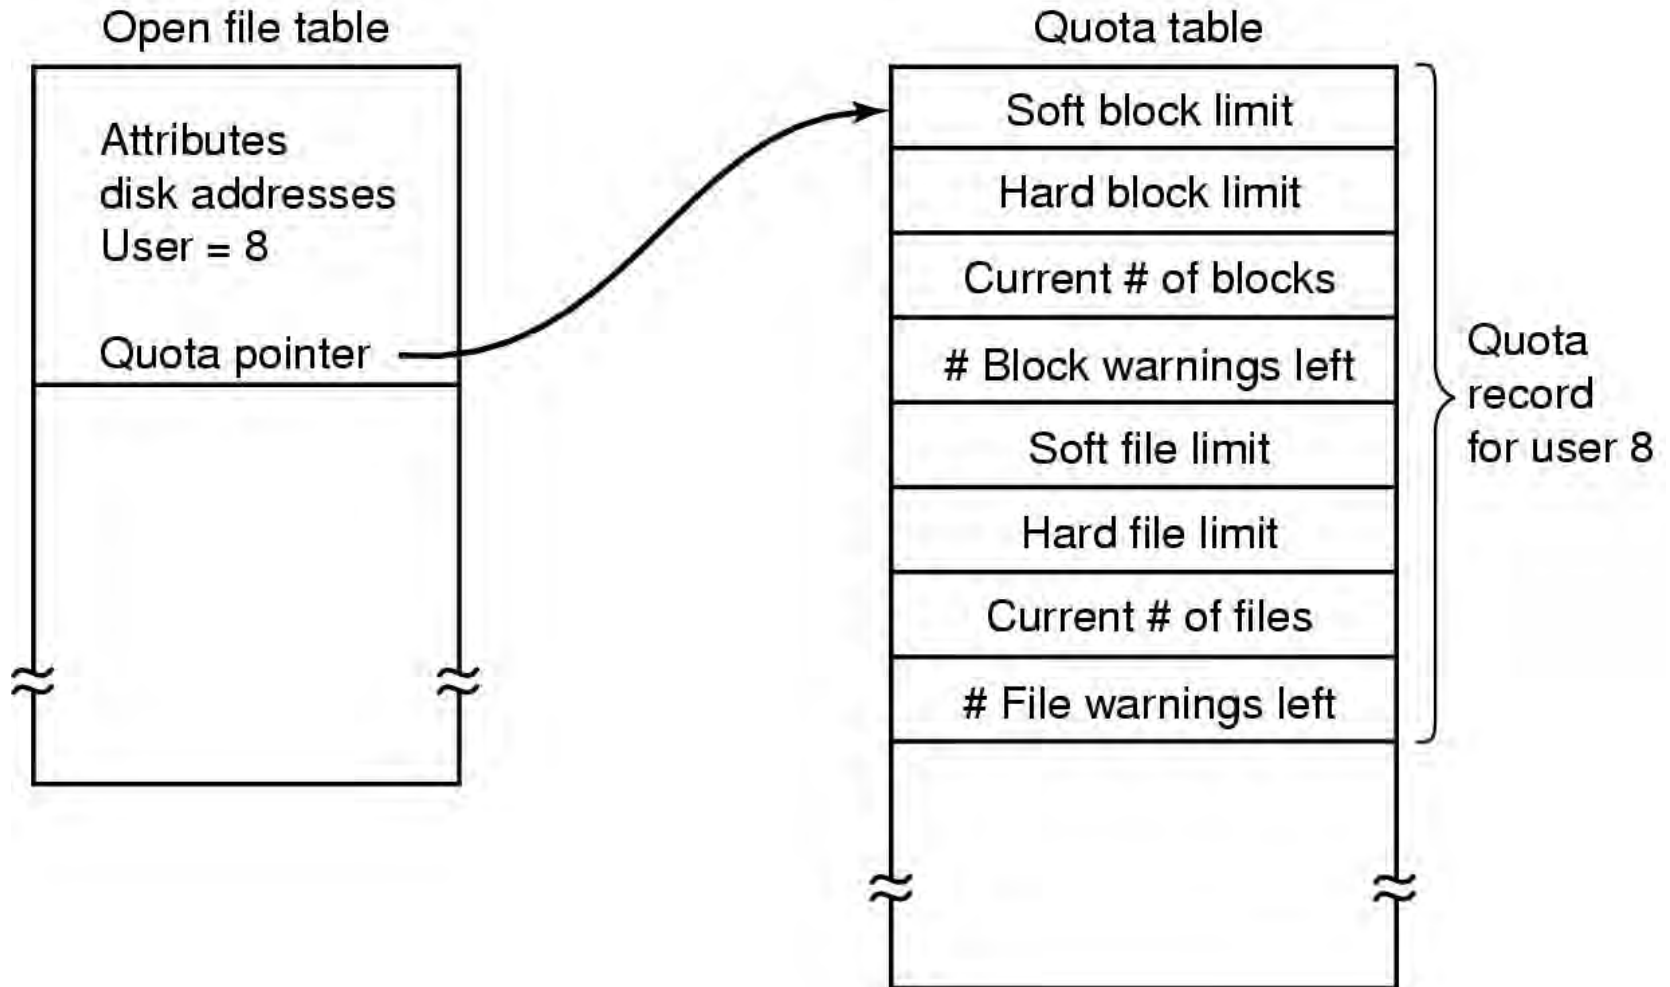


Keeping Track of Free Blocks



(a) Storing the free list on a linked list. (b) A bitmap.

Disk Quotas



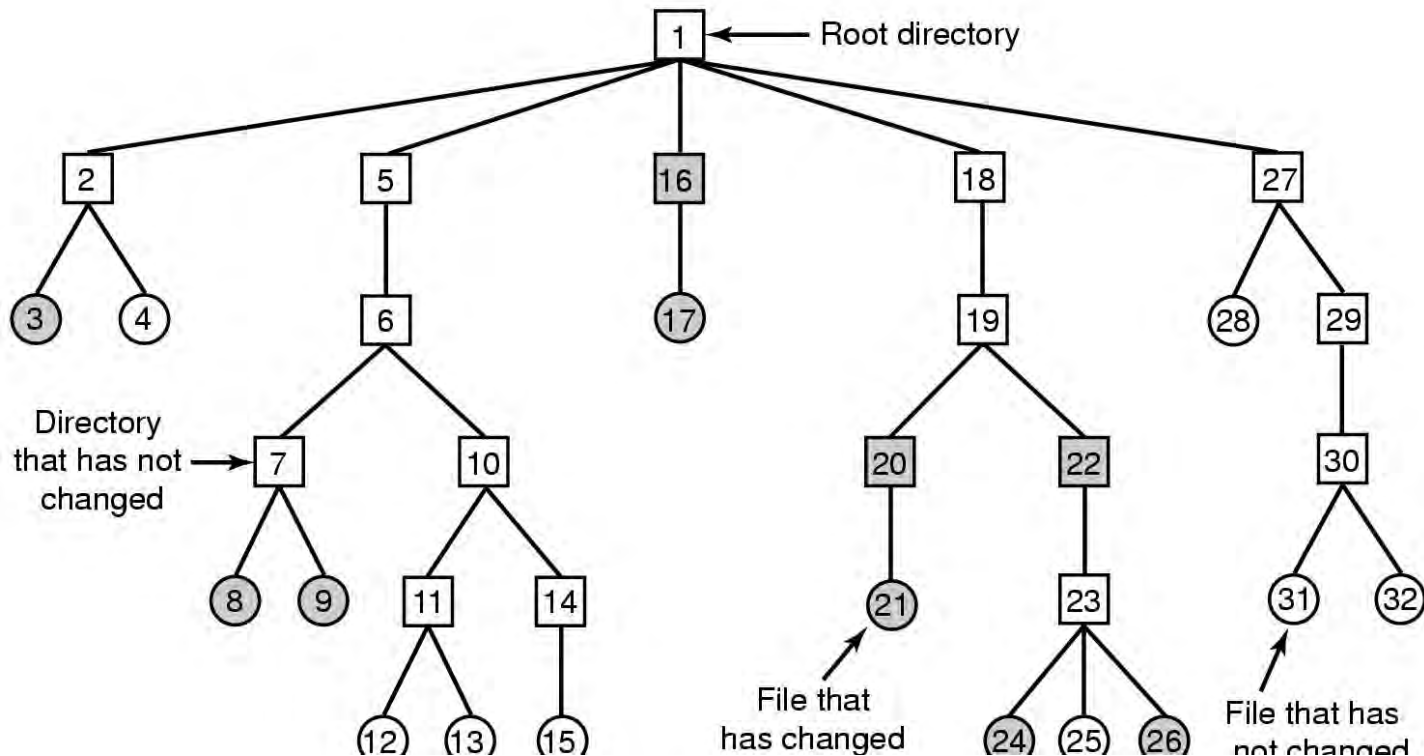


File System Backups

- Backup v.s. Recover
- Physical dump v.s. logical dump

File System Backups

- Full dump v.s. incremental dump



A file system to be dumped. Squares are directories, circles are files. Shaded items have been modified since last dump. Each directory and file is labeled by its i-node number.

File System Consistency

- Blocks in use
- Free blocks

Block number

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
1	1	0	1	0	1	1	1	1	0	0	1	1	1	0	0
Blocks in use															
0	0	1	0	1	0	0	0	0	1	1	0	0	0	1	1
Free blocks															

(a)

Block number

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
1	1	0	1	0	1	1	1	1	0	0	1	1	1	0	0
Blocks in use															
0	0	0	0	1	0	0	0	0	1	1	0	0	0	1	1
Free blocks															

(b)

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
1	1	0	1	0	1	1	1	1	0	0	1	1	1	0	0
Blocks in use															
0	0	1	0	2	0	0	0	0	1	1	0	0	0	1	1
Free blocks															

(c)

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
1	1	0	1	0	2	1	1	1	0	0	1	1	1	0	0
Blocks in use															
0	0	1	0	1	0	0	0	0	1	1	0	0	0	1	1
Free blocks															

(d)

(a) Consistent. (b) Missing block.

(c) Duplicate block in free list. (d) Duplicate data block




File System Performance

- Caching
 - Block cache, buffer cache
- I-nodes layout
- Block Read Ahead
- Reducing Disk Arm Motion


File System Performance: Cache

- Block cache
 - Buffer cache



File System Performance: Cache

- Block cache replacement algorithms
 - LRU, FIFO, ...



The buffer cache data structures



File System Performance: Cache


- Crash and File System Consistency?
- Write-back caches
- Write-through caches 通写高速缓存
- Sync, flushXXX




File System Performance

- Block Read Ahead
 - When executing a read from the disk, the disk arm moves the read/write head to (or near) the correct track, and after some settling time the read head begins to pick up bits.
 - Usually, the first sectors to be read are not the ones that have been requested by the operating system.
 - The disk's embedded computer typically saves these unrequested sectors in the disk buffer, in case the operating system requests them later.

File System Performance: I-nodes



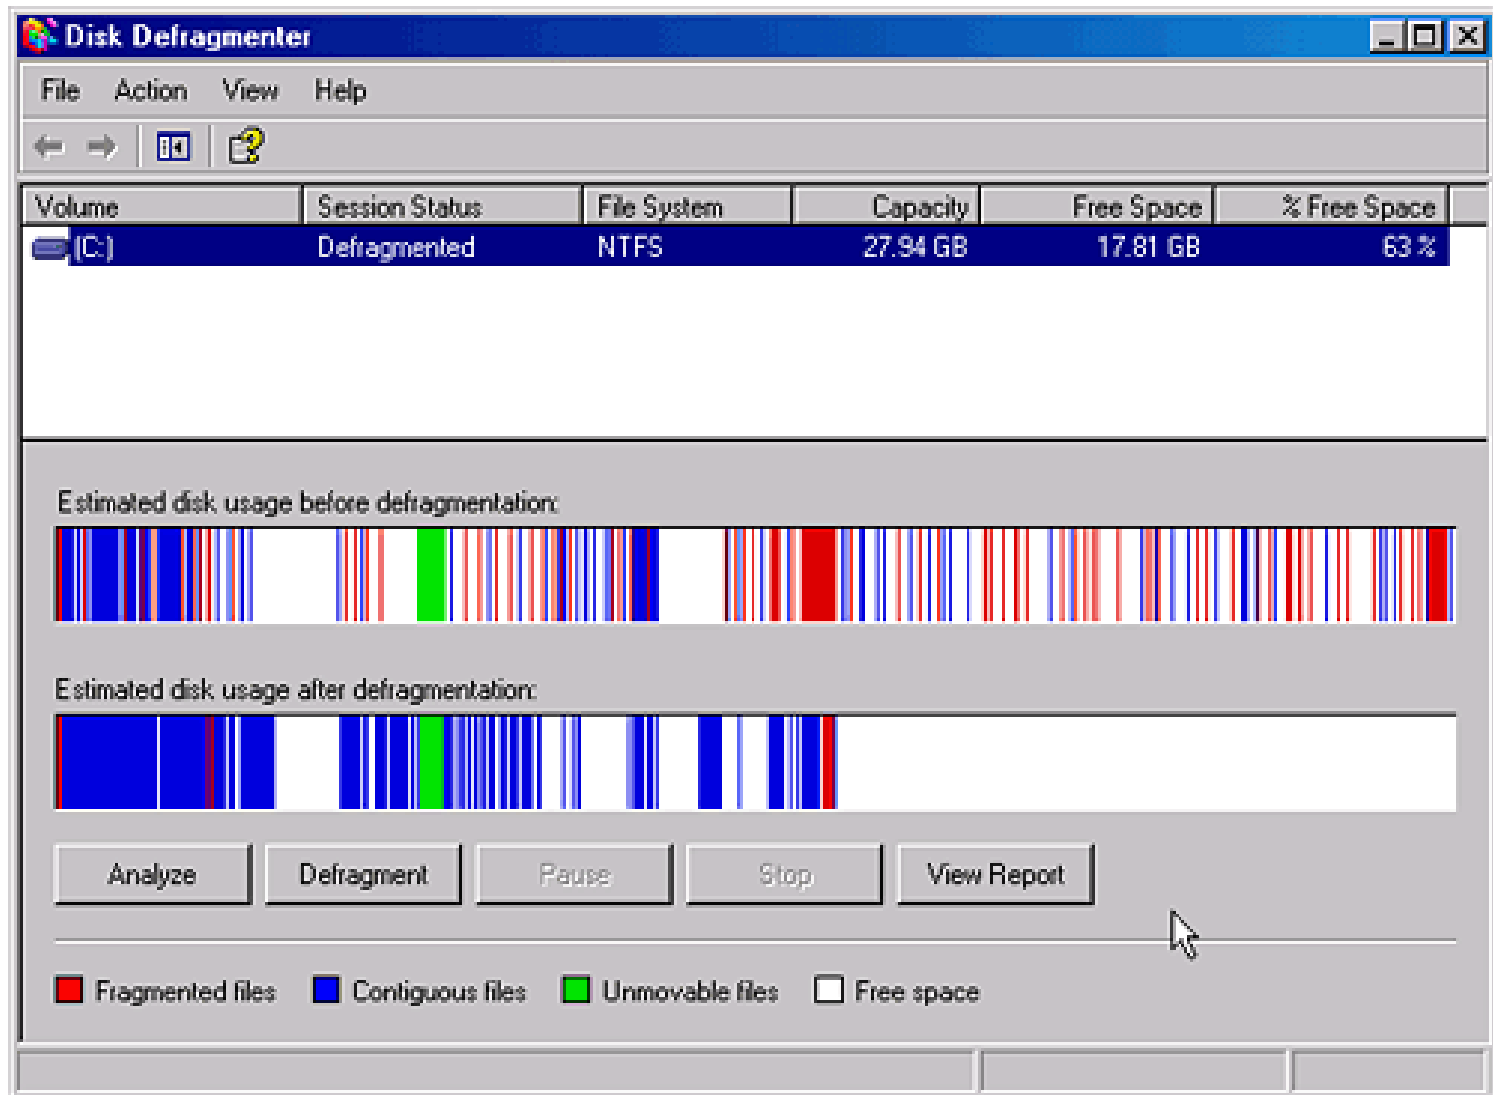
(a)



(b)

- **I-nodes** placed at the start of the disk
- Disk divided into cylinder groups
 - each with its own blocks and i-nodes

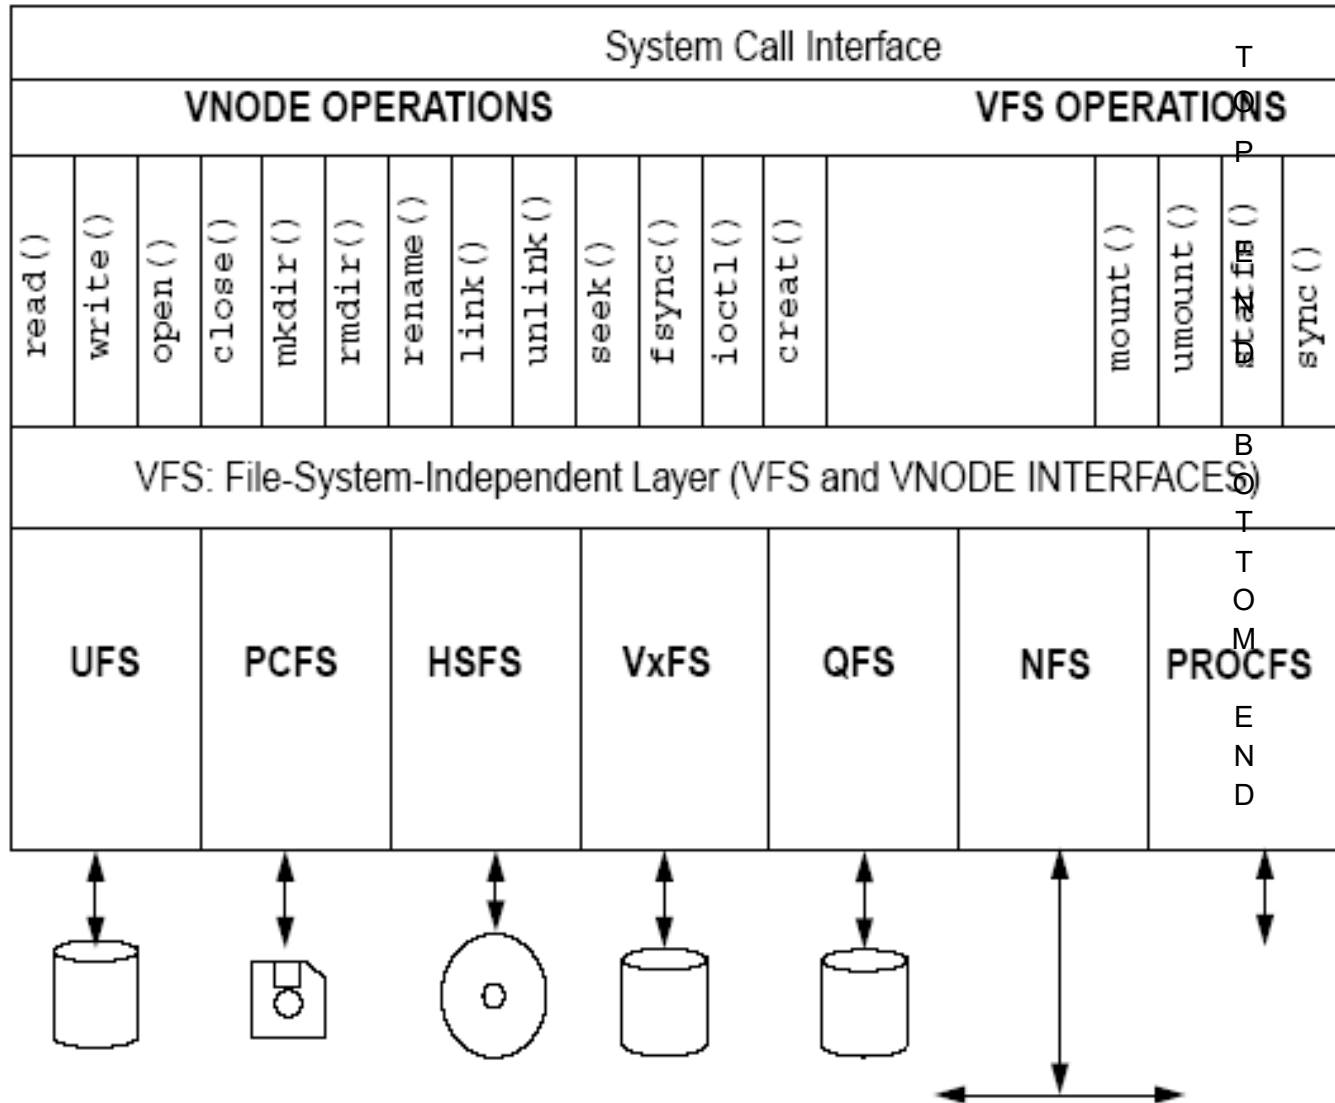
Defragmenting Disks



File System Framework Facilities

- **Loadable file system modules** are dynamically loaded at the time each file system type is first mounted.
- **The vnode/vfs framework** implements file functions and file system management functions.
- **File system caching** implements caching interface with the HAT layer of the virtual memory system to map, unmap, and manage the memory used for caching.
- **Path-name management** converts paths into vnode pointers.
- **Directory name caching** provides a mechanism to cache pathname-to-vnode mappings.

File System Layers





Topics of File Systems

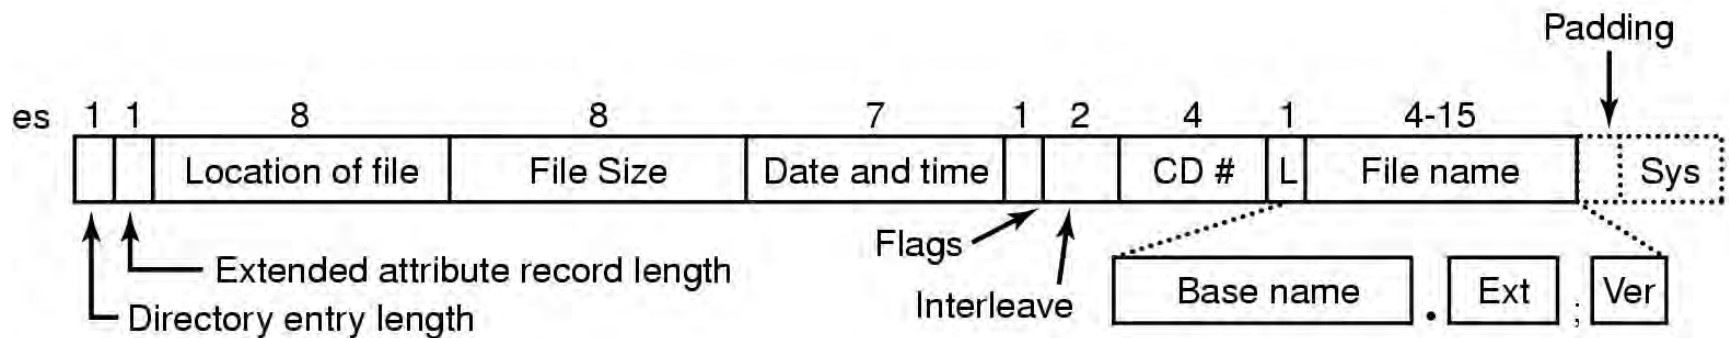
- Log-structured File System
- Journaling File System
- Network File System: NFS



File Systems Cases

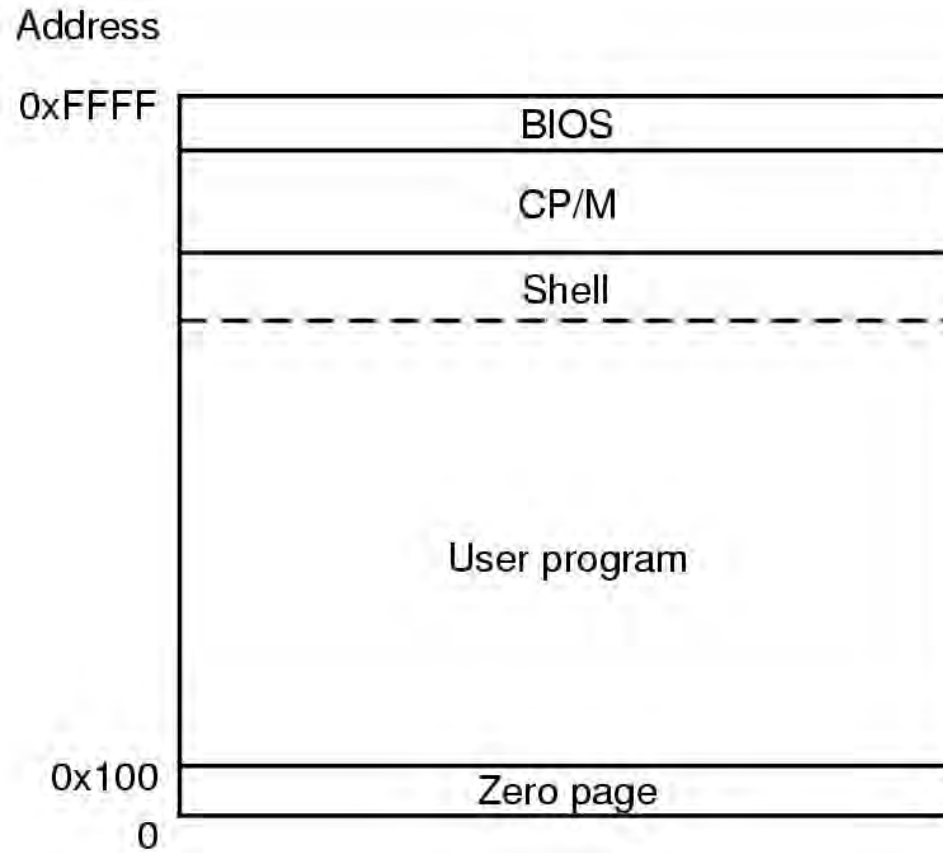
- CD-ROM: ISO9660
- CP/M File System
- FAT 16
- NTFS
- Ext4
- ...

CD-ROM File Systems




The ISO 9660 directory entry

The CP/M File System (1)




Memory layout of CP/M

The CP/M File System (2)



The CP/M directory entry format

The MS-DOS File System (1)



The MS-DOS directory entry

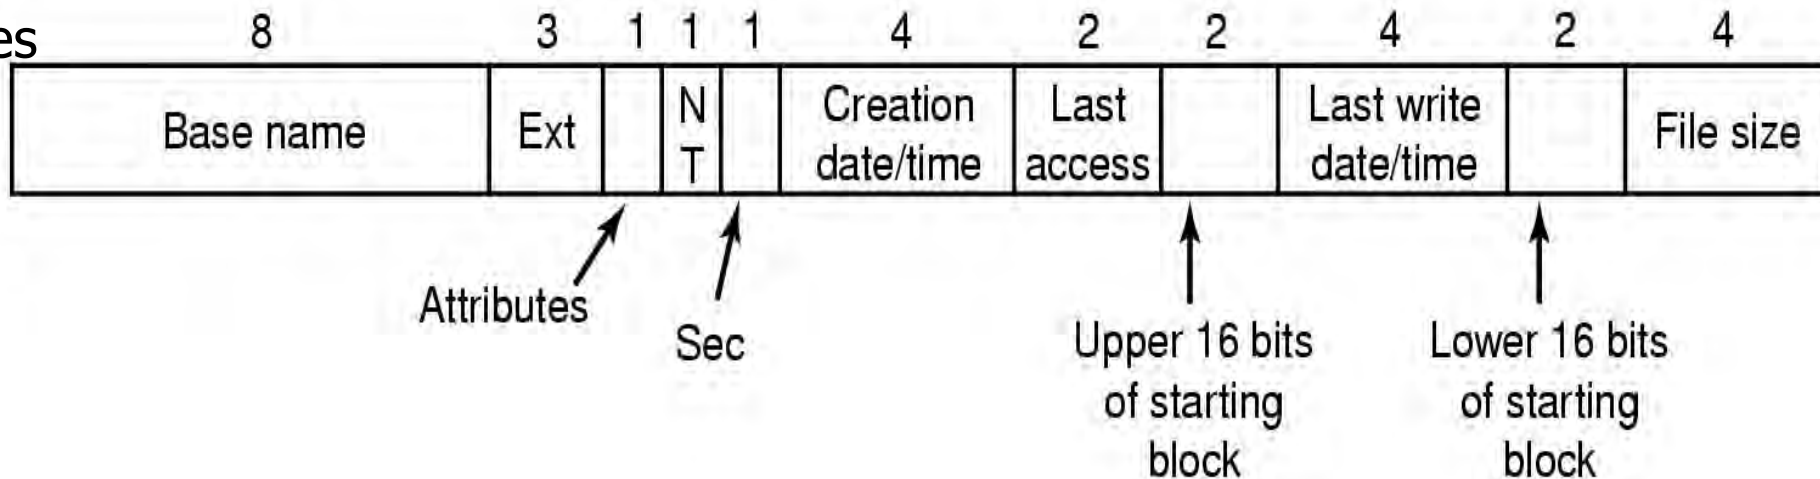
The MS-DOS File System (2)

Block size	FAT-12	FAT-16	FAT-32
0.5 KB	2 MB		
1 KB	4 MB		
2 KB	8 MB	128 MB	
4 KB	16 MB	256 MB	1 TB
8 KB		512 MB	2 TB
16 KB		1024 MB	2 TB
32 KB		2048 MB	2 TB

- Maximum partition for different block sizes
- The empty boxes represent forbidden combinations

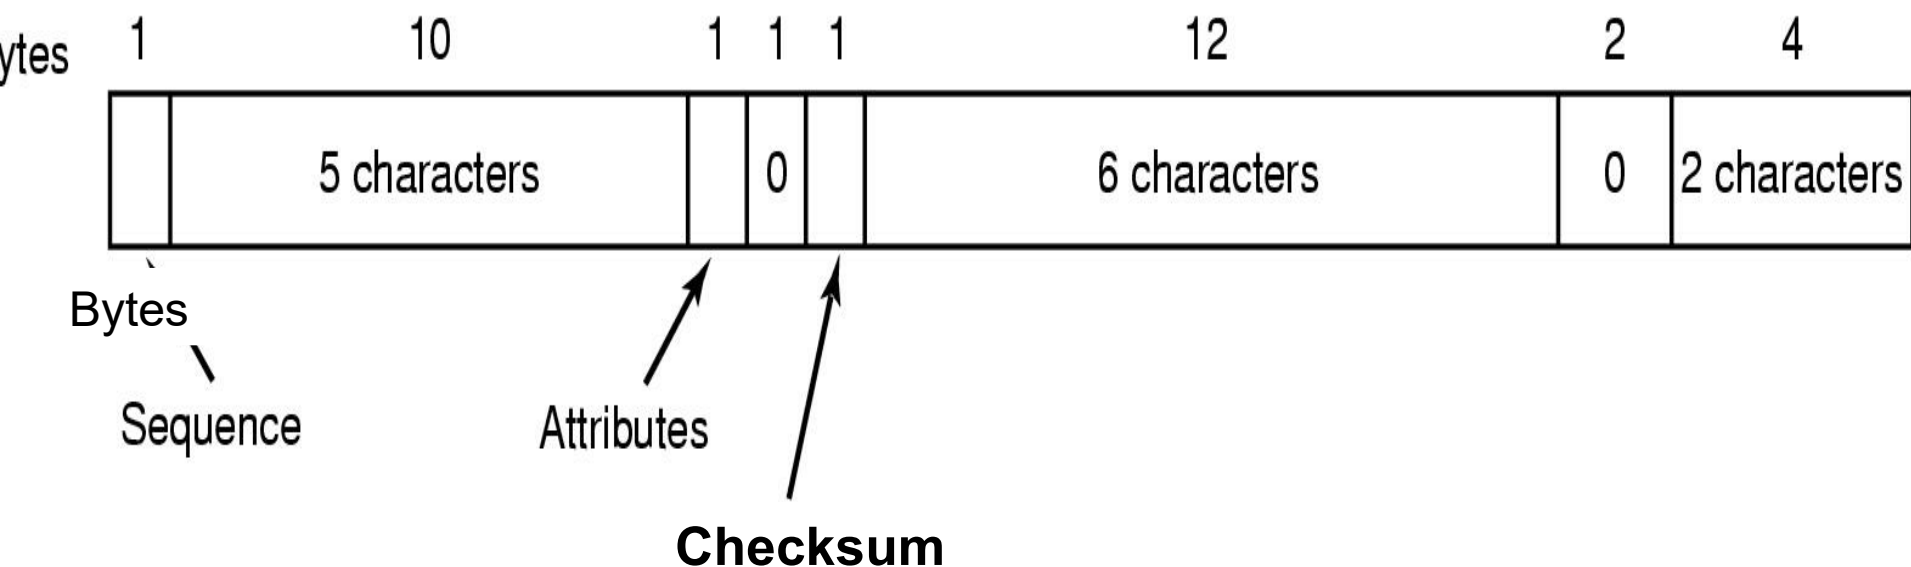
The Windows 98 File System 1/3

Bytes



The extended MOS-DOS directory entry used in Windows 98

The Windows 98 File System 2/3



An entry for (part of) a long file name in Windows 98

The Windows 98 File System 3/3

Bytes	68	d o g								A	0	C K					0					
	3	o v e								A	0	C K	t h e l a				0	z y				
	2	w n f o								A	0	C K	x j u m p				0	s				
	1	T h e q								A	0	C K	u i c k b				0	r o				
	T	H E Q U I ~ 1								A	N T	S	Creation time		Last acc	Upp	Last write		Low	Size		

An example of how a long name is stored in Windows 98



Summary

- File Concept
- Directory Concept
- File Share & Protection
- File System Implementation
- File System Reliability
- File System Performance
- File System Cases



Any Questions?